Trend research — 2026-04-19 · incr 20

one of the 12 trend judges was making things worse; a simple knob change fixes it

In two sentences

One of our 12 trend judges (**always_in**) was making things **worse**, not better — on 800 real trading days it gave an opinion **36%** of the time and was right **58%** of the time, which is **worse than just guessing 'UP' every day** (that baseline is 66%). Lengthening one of its knobs from 'look at the last 5 bars' to 'look at the last 10 bars' fixes it — now it gives an opinion on **52%** of days and is right **67%** of the time.

The whole story in a picture

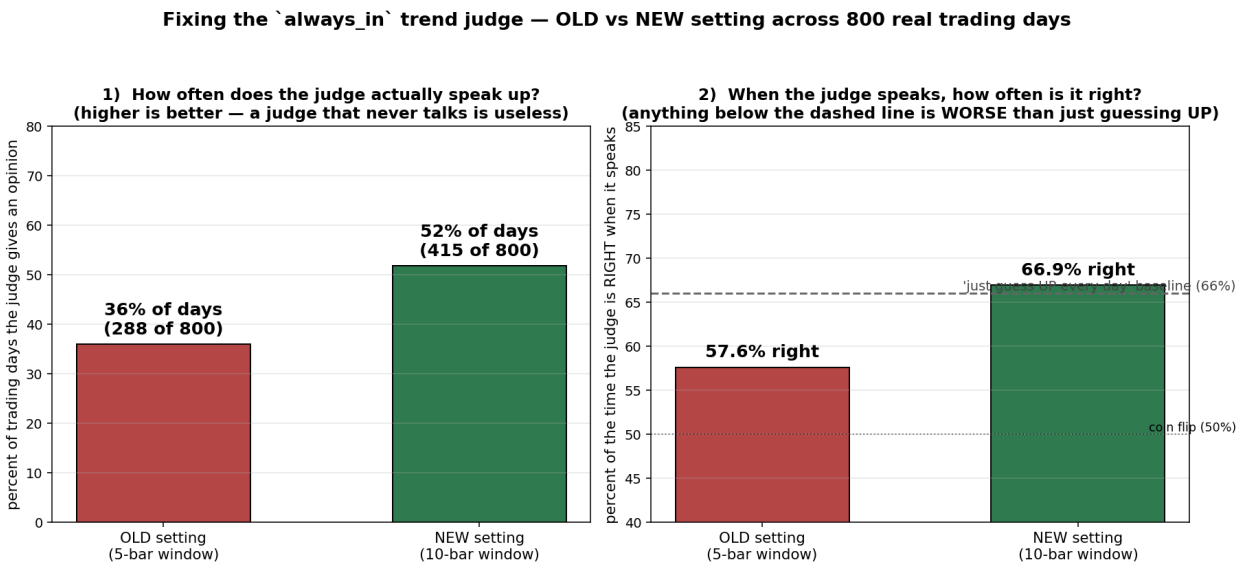

Two side-by-side questions. LEFT: how often does the judge speak up at all? RIGHT: when it speaks, how often is it right? The red bar is today's setting; the green bar is the proposed new setting. Both numbers get better — there's no trade-off.

Is the judge even worth listening to?

Three bars. The grey bar is the dumb strategy — every morning, predict 'up.' On this sample that's right **66%** of the time. The red bar is the judge at its current setting — **below** the grey bar, meaning the judge is *worse* than the dumb strategy. The green bar is the judge at the proposed new setting — **above** the grey bar, meaning it's actually adding real information.

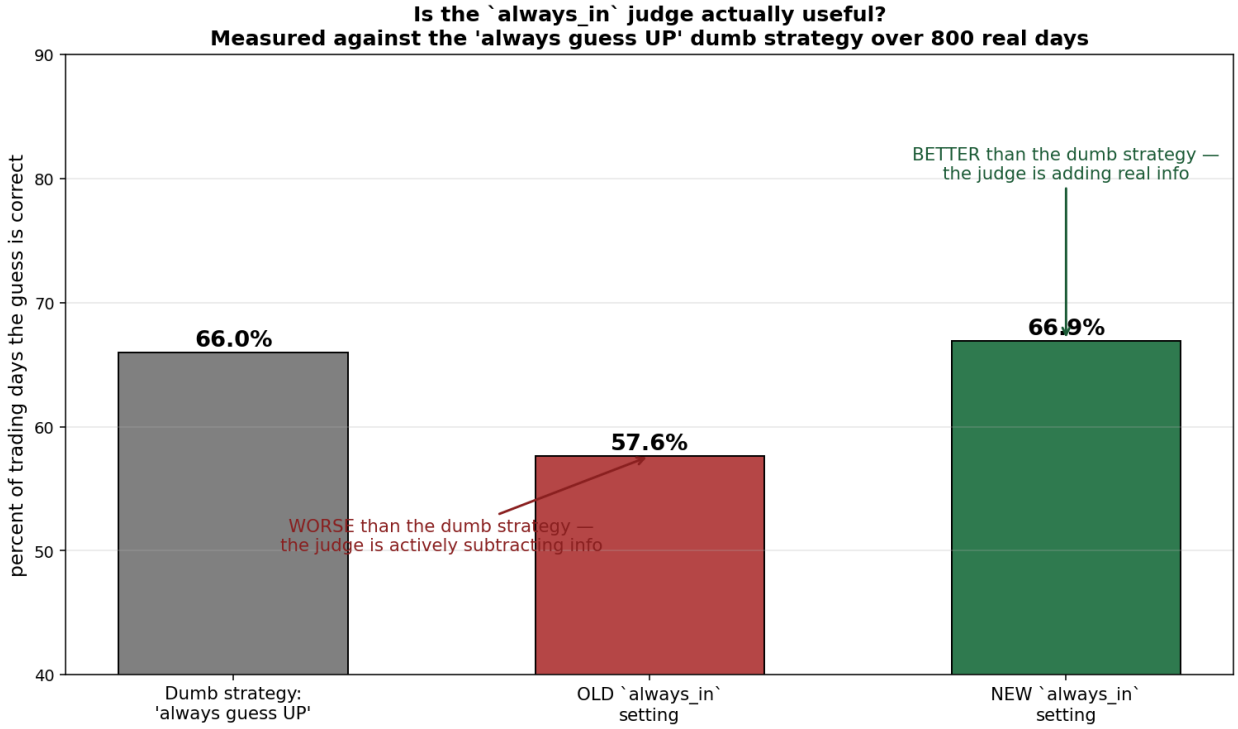

The red bar is worse than guessing. The green bar is better than guessing. That's the whole argument for the knob change.

What we did — no jargon

Our 'trend reading' is decided by a panel of 12 judges. Each judge has a couple of knobs that control how strict it is. One judge (**always_in**) looks at whether the market has been moving in one direction for a few bars in a row. It has two knobs:

- **'How many bars to look at'** — today it's set to **5**.
- **'How many bars in a row must agree'** — today it's set to **2**.

We tested every combination of 5 window lengths (3, 5, 7, 10, 15 bars) × 4 'in-a-row' values (1, 2, 3, 4) = **20 settings**, on **800 real trading days across 387 US stocks**. For each setting we asked two simple questions: (1) how often does the judge actually give an opinion? (2) when it does, how often is it right vs what the market actually did?

The result — on how many days does the judge speak?

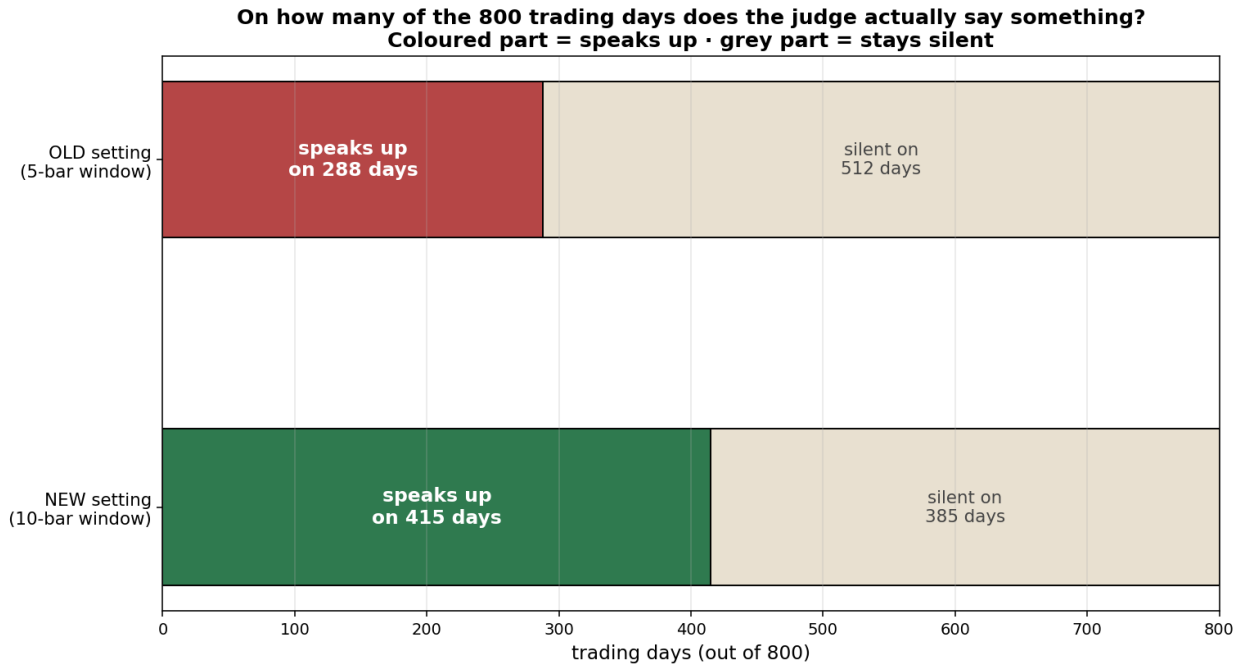

Coloured part = speaks up. Grey part = stays silent. The new setting doubles the size of the coloured part while also making its opinion more accurate.

What I'm recommending — still needs your nod

Change one number in the code: the bar-count window from **5** → **10**. Don't touch the 'in-a-row' knob. This is a read-only research pass — the number hasn't been changed, per the rule that thresholds need your explicit nod.

Caveat I have to name

'Right' here means 'the sign of the judge's opinion matches which way the day actually closed vs where it opened.' A judge that speaks LATE in the day has the easier job of agreeing with what already happened. So this is **fidelity-to-same-day**, not **predicts-the-future**. The eventual question — 'does this judge actually help the scanner make money?' — still needs the Pattern Lab win-rate database (still empty; still blocking).

For engineers — the 20-cell grid

Production (W=5, K=2) fires **36.0%**, acc **57.6%**. **Winner (W=10, K=2)** fires **51.9%**, acc **66.9%**. Binding knob is W, not K. K=1 is noise. High-conviction corner: W=15, K=1 hits 82.9% accuracy but only 5.5% fire rate.

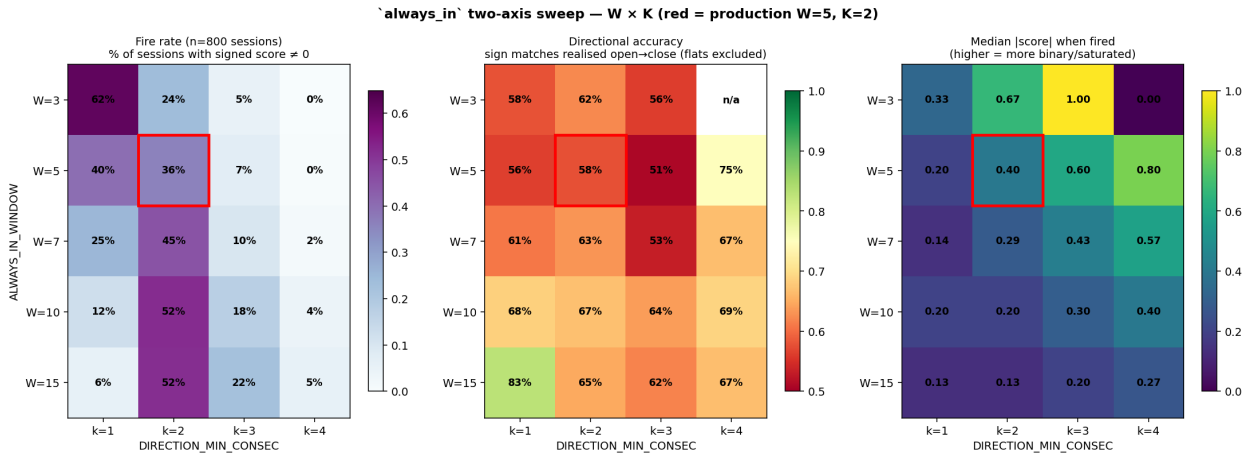

Red box = production. Left: fire rate. Middle: directional accuracy. Right: median |score| when fired.

Findings site: https://github.com/zerosumsystems-ui/iPhone-/tree/main/trend-classification